

# **SE 577**

## **Software Architecture**

### **The Software Architecture of Linux**

# Reference

I.T. Bowman, R.C. Holt, and N.V. Brewster: "*Linux as a Case Study Its Extracted Software Architecture*", Proc. ICSE '99, Los Angeles, CA., pp. 555-563.

# About the Paper

- Describes work done to document the software architecture of the Linux kernel
  - Linux kernel (~800KLOC) is responsible for:
    - Process management
    - Memory management
    - Hardware device management

# Why Study This Work?

- Look at an example of the architecture of a large, complex, software system.

# Why Study This Work?

- Look at an example of the architecture of a large, complex, software system.
- Observe the approach used to describe/document the architecture of the software

# Why Study This Work?

- Look at an example of the architecture of a large, complex, software system.
- Observe the approach used to describe/document the architecture of the software.
- Understand what needs to be done to figure out the architecture of such system [*extraction methodology*].

# Why Study This Work?

- Look at an example of the architecture of a large, complex, software system.
- Observe the approach used to describe/document the architecture of such software system.
- Understand what needs to be done to figure out the architecture of such system [extraction methodology].
- Point out the differences and significance between the **conceptual** and **concrete** architectures?

# Architectural Documentation

- Two types of architectural documentation:
  - Conceptual architecture  
(*Prescriptive* architecture)
    - How developers think about the system
    - Relationships between subsystems that are meaningful to developers
  - Concrete architecture  
(*Descriptive* architecture)
    - Relationships that exist in the implemented system.



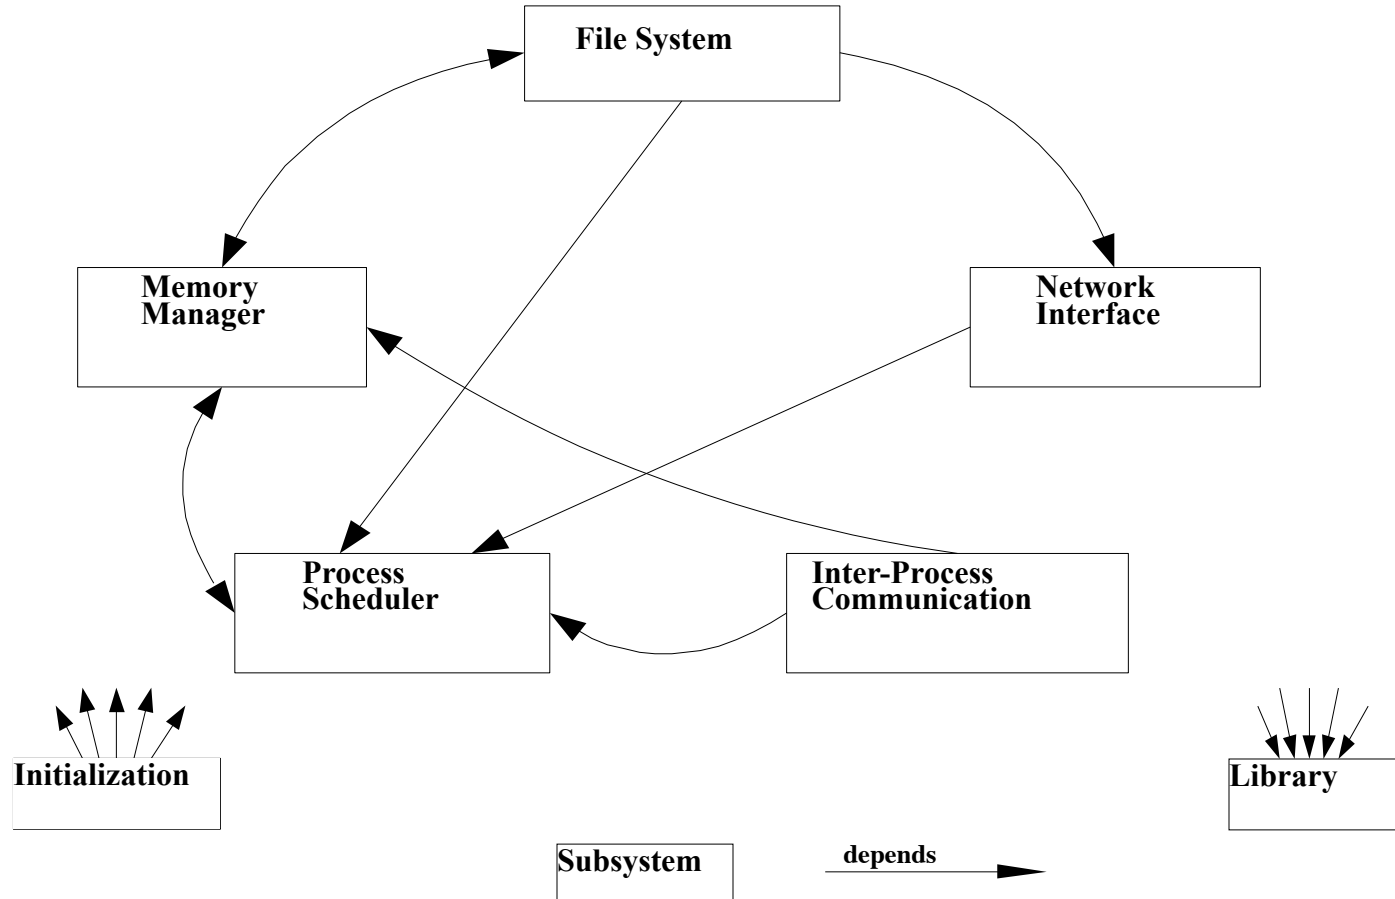
# Architectural Documentation (cont'd)

- Architecture documentation
  - Individual functions and even modules are not described in detail.
  - Subsystems and relations between them are documented.
- Architectural documentation is valuable for understanding the system.
- Great help for evaluation or re-engineering efforts.

# Architectural Documentation (cont'd)

- Unless architectural documentation is maintained it will become obsolete as the system undergoes further changes.
- Others have proposed work to keep architectural documentation up to date (e.g., *Software Bookshelf* by Finnigan, et al.)

# Conceptual Architecture



# Conceptual Architecture (cont'd)

- Diagram shows “depends-on” relationship
  - E.g., Memory Manager depends on the File System to swap memory to and from disk.
- Each subsystem has additional subsystems hierarchically nested within it (not shown).
- The Initialization subsystem depends on all other kernel subsystems (it calls initialization routines).
- All the kernel subsystems depend on the Library subsystem.

# Conceptual Architecture (cont'd)

Seven major subsystems:

- **Process Scheduler**
  - Responsible for supporting multitasking by changing which user process executes.
- **Memory Manager**
  - Provides a separate memory space for each user process and uses swapping to support more processes than fit in physical memory.
- **File System**
  - Provides access to hardware devices. User processes can access keyboards, tape drives, hard drives, and modems using one interface that is implemented by the File System.

# Conceptual Architecture (cont'd)

- Network Interface
  - Encapsulates access to network devices in a similar manner to the File System. User processes can communicate with other computers using several different types of network hardware and transmission protocols.
- InterProcess Communication
  - The IPC subsystem allows user processes to communicate with other processes on the same Computer. Synchronization memory sharing and interprocess messaging primitives are supported by the IPC subsystem.

# Conceptual Architecture (cont'd)

- Initialization
  - Responsible for initializing the rest of the Linux kernel with appropriate user configured settings.
- Library
  - Contains routines which are used throughout the kernel.

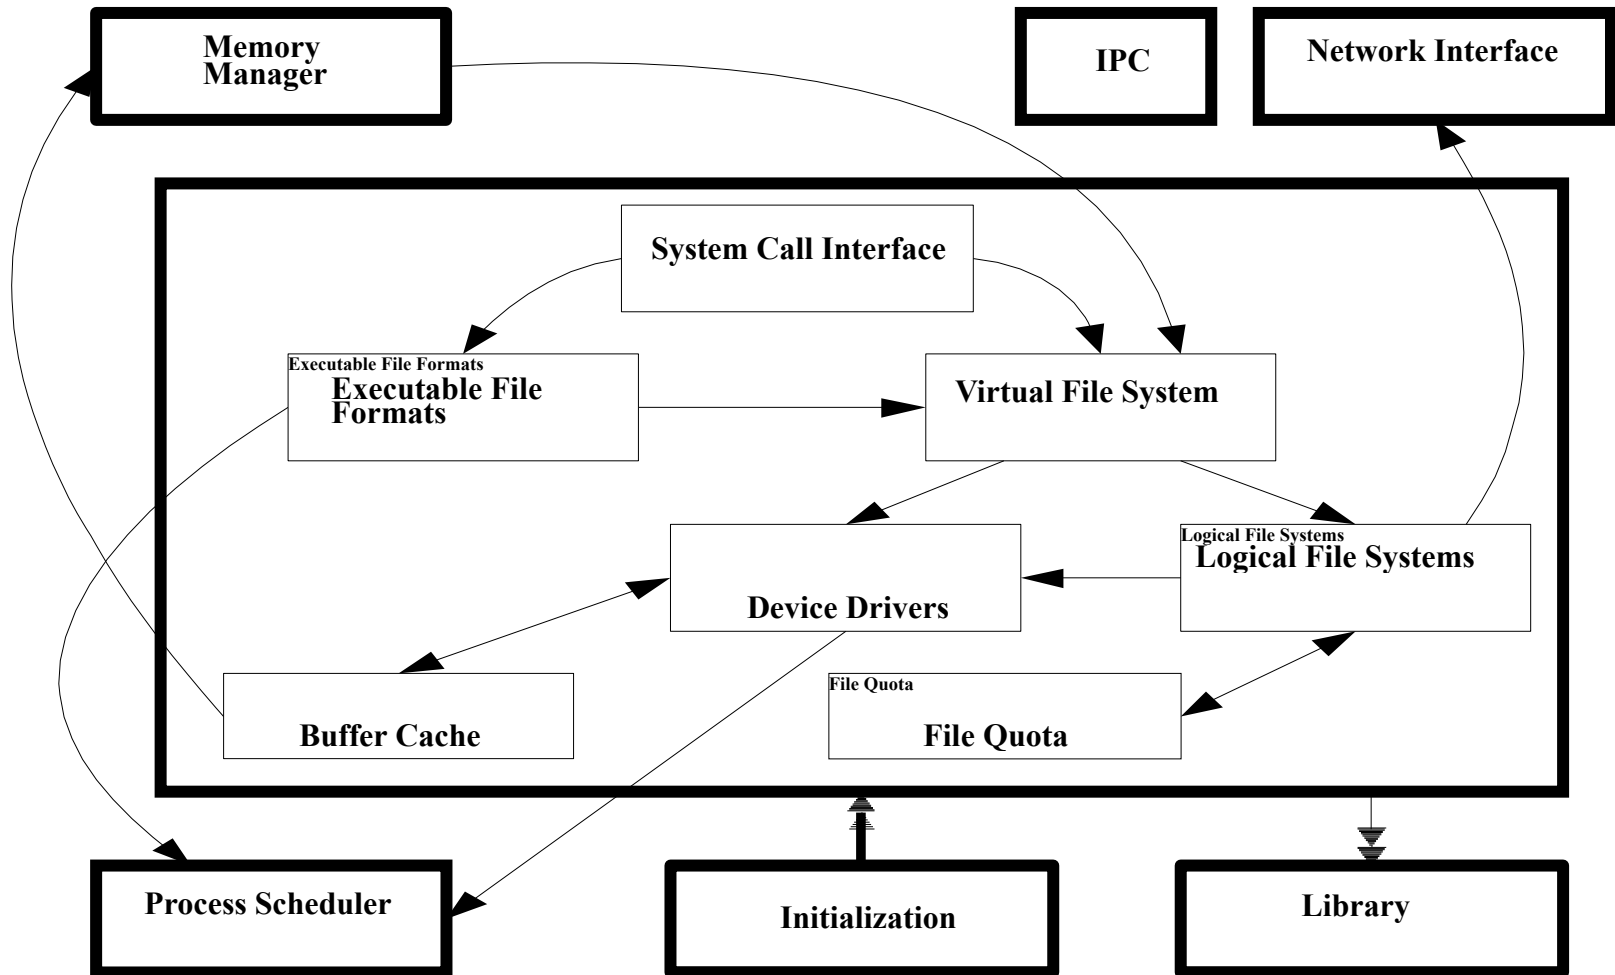
# File System Conceptual Architecture

Three main roles:

- Provide access to a wide variety of hardware devices.
- Support several different logical file system formats that control how files are mapped to physical locations on hardware devices.
- Allow programs to be stored in several executable formats including interpreted scripts.



# File System Conceptual Architecture (cont'd)



# The File System

Five subsystems implement the main roles of the File System:

- **Device Drivers**
  - Perform all communication with hardware devices supported by Linux.
- **Logical File Systems**
  - Implement several logical file systems that can be placed on hardware devices; these different file systems allow interoperability with different operating systems and also allow specialized functionality such as encryption compression and high performance.

# The File System (cont'd)

- Executable File Formats
  - Allow clients to execute programs from several different executable file formats including not only compiled programs but also interpreted scripts.
- File Quota
  - Allow system administrators to limit the amount of file storage that individual users may use.
- Buffer Cache
  - Provide memory buffers for input/output operations and reduces hardware accesses by caching data and eliminating redundant reads and writes.

# The File System (cont'd)

Unexpected relationships from Conceptual architecture:

- **Device Drivers → Process Scheduler**
  - A device driver might request to be suspended while waiting for i/o to complete.
- **Logical File System → Network Interface**
  - Some logical files are stored on another computer and accessed using the network.
- **Memory Manager → Virtual File System**
  - Swap memory to and from secondary storage.
  - [Not sure why this was unexpected...]

# The File System (cont'd)

- Linux uses the Façade design pattern to allow user processes and other parts of the kernel to use elements of the File System through a single interface.
- The File System follows the “object-oriented” or “data abstraction” style (described by Garlan and Shaw).

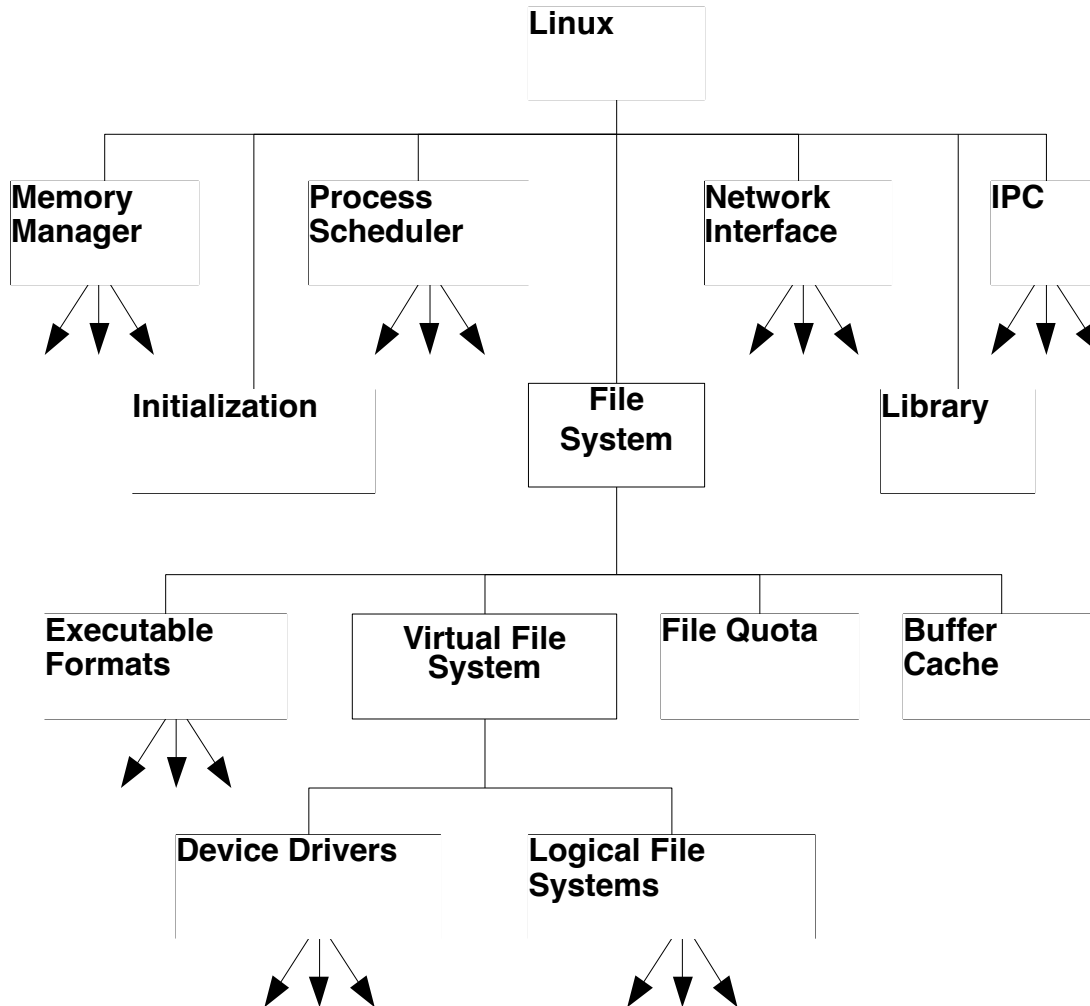
# Extracting the Concrete Architecture

- Kernel contains 1,682 source files.
- Relations between source files are at too low a level for easy system understanding.
- Focus on relations between subsystems.
- Used tools (grok and lsdedit) to determine what relations exist between subsystems, based on the relations between the source files that are contained in the subsystems.
- Combined tool support and human interpretation to extract the concrete architecture.

# Extraction Methodology


- Manually created a tree structure of subsystems.
- Assigned each subsystem to a single containing subsystem (used the subsystems from conceptual architecture.)
- Manually assigned source files to subsystems based on:
  - directory structure
  - file naming conventions
  - source code comments
  - documentation
  - examination of the source code
- If a set of source files seemed logically related, created a new subsystem to contain them.

# Hierarchical Decomposition (Partial)






# Concrete Architecture



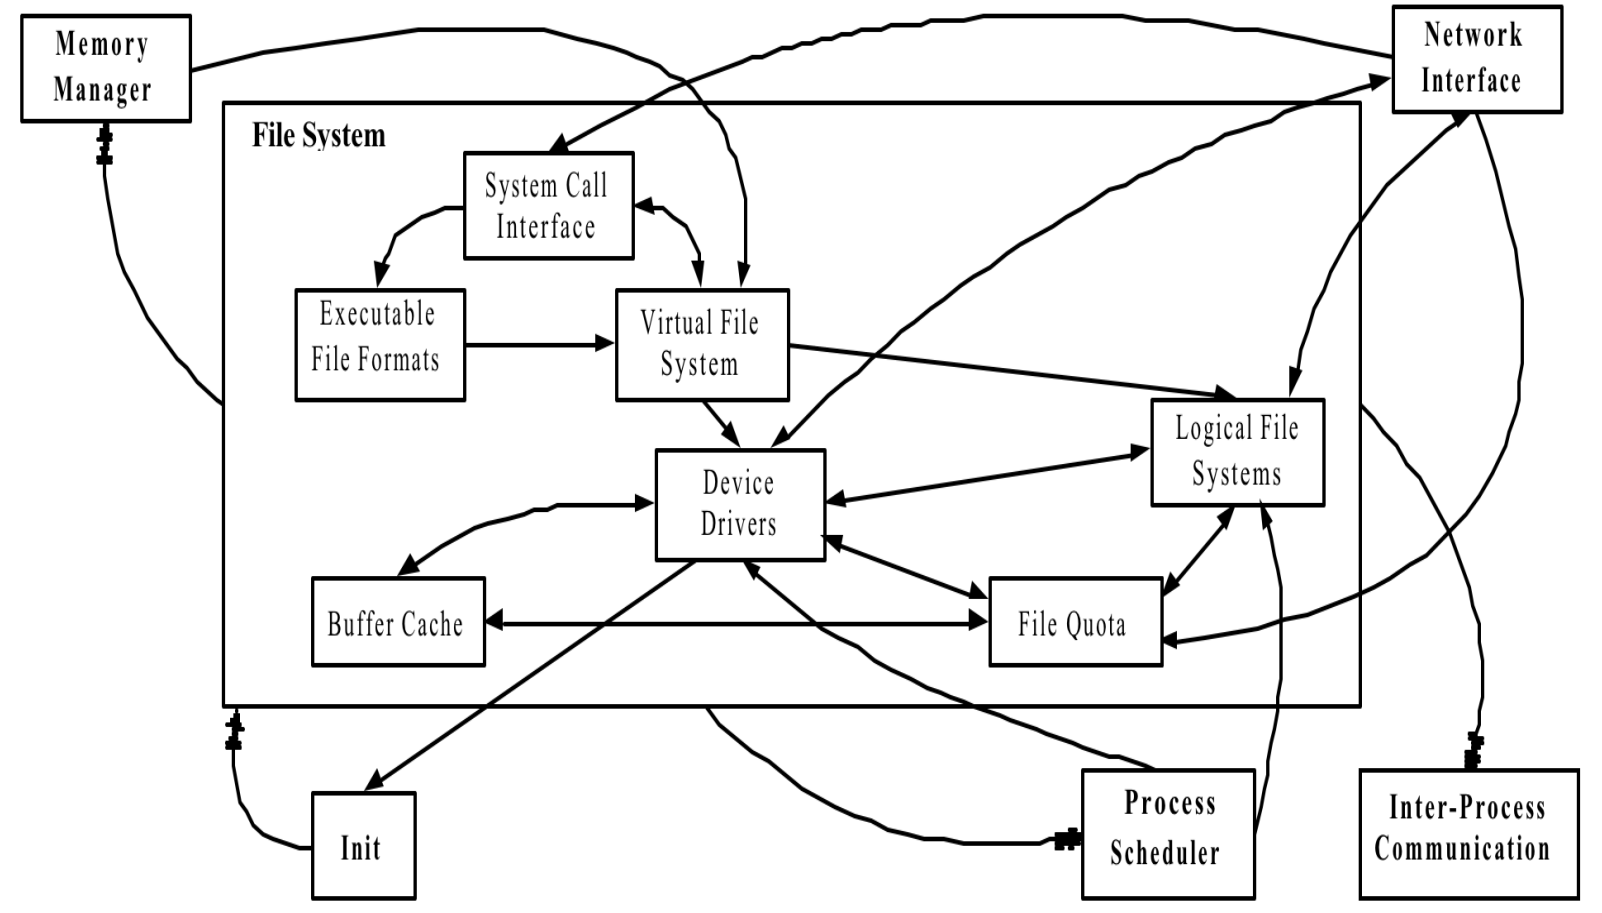
# Conceptual Architecture



# Concrete vs Conceptual Architecture

- Same subsystems
- Dependency relations appear to be quite different
  - Concrete architecture almost fully connected.
- Lesson learned: a concrete implementation is likely to have more dependencies. Why?
  - Developers might avoid existing interfaces for better efficiency.
  - Unclear if the dependencies are required.

# File System Concrete Architecture



# File System Concrete Architecture (cont'd)

- Same subsystems.
- Also has substantially more dependency relations. Why?
  - Missed dependencies in the documentation.
  - Functionality that was implemented in multiple subsystems.
  - Unexpected control flow implementations.

# File System Concrete Architecture (cont'd)

- Some unexpected relations:
  - Network Interface → File System
    - The Network Interface directly calls functions in the implementation of two logical file systems (NCPFS, SMBFS).
  - Process Scheduler → Device Drivers
    - A Process Scheduler routine calls a routine implemented within the Device Drivers.
  - File System subsystems → IPC
    - Synchronization primitives are also used by the kernel.

# Logical File System Concrete Architecture

- Contains 17 different logical file systems.
- They map logical files to physical locations on storage devices.
- Has more dependencies than predicted in the conceptual architecture.

# Takeaways

- Automated tools are very important, but tools alone are not enough.
- Human participation still required.
  - Source code examination is part of the deal
- Remember: the architecture of a system is not “visible”.
- Substantial differences between conceptual vs. concrete architecture
  - Missed the use of subsystems (e.g., IPC implements mechanisms used throughout the kernel).
  - Existing interfaces are bypassed for efficiency reasons.



## Takeaways (cont'd)

- The concrete architecture should be used to refine the conceptual architecture.
- May not be desirable to add all relations
  - Many are not essential and hinder system understanding.
- Concrete architecture identifies potential improvements.